

Comparaison des Modèles ARIMA, SARIMA et LSTM pour la Prédiction des Prix de Clôture des Actions

Par: Mr HACHIMI ATMANE



pandas



NumPy



scikit
learn

```
In [359... # Importer les bibliothèques nécessaires
import yfinance as yf
from datetime import datetime
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.preprocessing import MinMaxScaler
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.stattools import adfuller
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from statsmodels.tsa.arima.model import ARIMA
from statsmodels.tsa.statespace.sarimax import SARIMAX
from keras.models import Sequential
from keras.layers import Dense, LSTM
from sklearn.metrics import mean_squared_error, mean_absolute_error
```

```
In [360... #download data from yfinance
end = datetime.now()
start = datetime(end.year - 3, end.month, end.day)

data = yf.download('AAPL', start=start, end=end)
```

```
[*****100%*****] 1 of 1 completed
```

```
In [361... #display the data
data.head()
```

Out[361...

	Open	High	Low	Close	Adj Close	Volume
Date						
2021-12-14	175.250000	177.740005	172.210007	174.330002	171.520370	139380400
2021-12-15	175.110001	179.500000	172.309998	179.300003	176.410278	131063300
2021-12-16	179.279999	181.139999	170.750000	172.259995	169.483704	150185800
2021-12-17	169.929993	173.470001	169.690002	171.139999	168.381775	195432700
2021-12-20	168.279999	170.580002	167.460007	169.750000	167.014175	107499100

In [362...

```
#display data infos
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 755 entries, 2021-12-14 to 2024-12-13
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   Open        755 non-null    float64
 1   High        755 non-null    float64
 2   Low         755 non-null    float64
 3   Close       755 non-null    float64
 4   Adj Close   755 non-null    float64
 5   Volume      755 non-null    int64
dtypes: float64(5), int64(1)
memory usage: 41.3 KB
```

In [363...

```
#verify the index of the data
data.index
```


Out[363...

```
DatetimeIndex(['2021-12-14', '2021-12-15', '2021-12-16', '2021-12-17',
               '2021-12-20', '2021-12-21', '2021-12-22', '2021-12-23',
               '2021-12-27', '2021-12-28',
               ...,
               '2024-12-02', '2024-12-03', '2024-12-04', '2024-12-05',
               '2024-12-06', '2024-12-09', '2024-12-10', '2024-12-11',
               '2024-12-12', '2024-12-13'],
              dtype='datetime64[ns]', name='Date', length=755, freq=None)
```

In [364...

```
#plot evry column vs time

data.plot(subplots=True)
plt.show()
```



```
In [365... # Check for missing values in the dataset
data.isnull().sum()
```

```
Out[365... Open      0
High      0
Low       0
Close     0
Adj Close  0
Volume    0
dtype: int64
```

```
In [366... # Count the number of duplicate rows in the DataFrame and return the count
data.duplicated().sum()
```

```
Out[366... 0
```


```
In [367... #create a new dataframe with the one column closing price
df = data['Close'].to_frame()
df.head()
```

Out[367...

Close

Date	
2021-12-14	174.330002
2021-12-15	179.300003
2021-12-16	172.259995
2021-12-17	171.139999
2021-12-20	169.750000

```
In [ ]: #use the seasnal_decompose function from statsmodels to decompose
# the closing price into trend, seasonal and residual components
from statsmodels.tsa.seasonal import seasonal_decompose
result = seasonal_decompose(df, model='multiplicative', period=30)
result.plot()
plt.show()
```



Description du graphique de décomposition saisonnière

Le graphique de décomposition saisonnière de la série temporelle des prix de clôture se compose de quatre sous-graphiques :

1. **Observed (Observé) :**

- Le sous-graphe "Observed" montre les valeurs originales des prix de clôture sur la période spécifiée. On peut observer les fluctuations quotidiennes des prix de clôture.

2. **Trend (Tendance) :**

- Le sous-graphe "Trend" montre la tendance à long terme des prix de clôture. On peut voir une tendance générale à la hausse des prix de clôture sur la période. Cette composante lisse les fluctuations à court terme pour révéler la direction générale de la série.

3. **Seasonal (Saisonnalité) :**

- Le sous-graphe "Seasonal" montre les variations saisonnières répétitives des prix de clôture. On peut observer des motifs périodiques qui se répètent à intervalles réguliers. Cette composante capture les motifs saisonniers dans les prix de clôture.

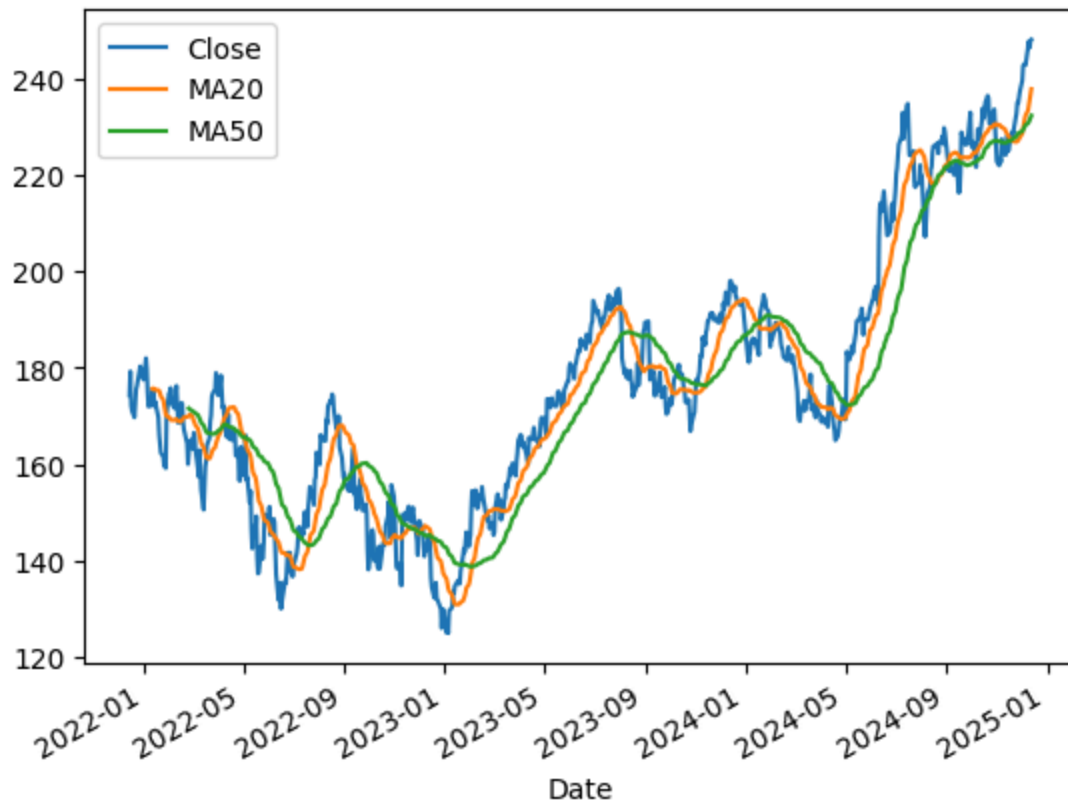
4. **Residual (Résidu) :**

- Le sous-graphe "Residual" montre les résidus ou les erreurs après avoir retiré la tendance et la saisonnalité. On peut voir les variations aléatoires ou les bruits qui ne sont pas expliqués par la tendance ou la saisonnalité. Cette composante représente les variations imprévisibles dans les prix de clôture.

En résumé, la décomposition saisonnière permet de séparer la série temporelle en ses composantes de tendance, de saisonnalité et de résidu, ce qui facilite l'analyse et la compréhension des motifs sous-jacents dans les prix de clôture.

In [369...


```
#the moving average of the closing price over a 20-day and 50-day period for visual
df['MA20'] = df['Close'].rolling(window=20).mean()
df['MA50'] = df['Close'].rolling(window=50).mean()
df[['Close', 'MA20', 'MA50']].plot()
plt.show()
```



Calculer les **moyennes mobiles** sur 20 et 50 jours permet de lisser les fluctuations quotidiennes des prix de clôture des actions, facilitant ainsi l'identification des tendances à court et à long terme. Cela aide les investisseurs à prendre des décisions éclairées en visualisant les mouvements de prix plus stables et en détectant les signaux de trading potentiels.

In [370...

```
#Scale the closing price using MinMaxScaler for better visualization and analysis
scaler = MinMaxScaler()
df['Close_scaled'] = scaler.fit_transform(df[['Close']])
df['Close_scaled'].plot()
plt.show()
```



Scaling des données est une étape cruciale pour assurer que les caractéristiques sont sur une échelle commune, améliorer les performances des algorithmes de machine learning, éviter les problèmes de stabilité numérique et faciliter la visualisation des données.

```
In [371... #display the scaled closing price
df['Close_scaled'].head()
```

```
Out[371... Date
2021-12-14    0.400536
2021-12-15    0.440907
2021-12-16    0.383722
2021-12-17    0.374624
2021-12-20    0.363334
Name: Close_scaled, dtype: float64
```

```
In [372... # Define a function to test the stationarity of a time series using the Augmented D
def test_stationarity(timeseries):

    #Determing rolling statistics
    rolmean = timeseries.rolling(12).mean()
    rolstd = timeseries.rolling(12).std()

    #Plot rolling statistics:
    plt.plot(timeseries, color='blue',label='Original')
    plt.plot(rolmean, color='red', label='Rolling Mean')
    plt.plot(rolstd, color='black', label = 'Rolling Std')
    plt.legend(loc='best')
    plt.title('Rolling Mean and Standard Deviation')
    plt.show(block=False)
```

```


print("Results of dickey fuller test")
adft = adfuller(timeseries,autolag='AIC')

# output for dft will give us without defining what the values are.
#hence we manually write what values does it explains using a for loop

output = pd.Series(adft[0:4],index=['Test Statistics','p-value','No. of lags us
for key,values in adft[4].items():
    output['critical value (%s)'%key] = values
print(output)

test_stationarity(df['Close'])

```



```

Results of dickey fuller test
Test Statistics          -0.271963
p-value                 0.929382
No. of lags used        0.000000
Number of observations used  754.000000
critical value (1%)      -3.439053
critical value (5%)      -2.865381
critical value (10%)     -2.568815
dtype: float64

```



In [373...

```

#plot the autocorrelation and partial autocorrelation functions
plot_acf(df['Close'], lags=20)
plt.show()

plot_pacf(df['Close'], lags=20)
plt.show()

```

Analyse des graphes ACF et PACF

1-ACF (Autocorrelation Function)

Observation : Les coefficients sont significatifs (en dehors de la zone bleue) pour de nombreux retards (lags) au début, et ils décroissent lentement au fur et à mesure.

Interprétation : Cela indique la présence d'une dépendance à long terme dans les données.

La décroissance lente suggère une possible non-stationnarité dans les données ou une série avec une tendance. Conclusion : Une différenciation peut être nécessaire pour rendre la série stationnaire, un prérequis pour les modèles ARIMA.

2- PACF (Partial Autocorrelation Function)

Observation : Seuls les premiers retards (1 ou 2) sont significatifs ; les suivants sont proches de zéro. Interprétation : Cela indique que la dépendance est principalement capturée par le premier ou les deux premiers termes d'autoregression. Conclusion : Cela suggère qu'un modèle AR(1) ou AR(2) pourrait être suffisant pour expliquer les dépendances dans les données. Conclusions générales Les données montrent une non-stationnarité potentielle.

Une différenciation est probablement nécessaire. La structure PACF suggère qu'un modèle AR(1) ou AR(2) peut être un bon point de départ. Si la différenciation est effectuée (modèle ARIMA), le paramètre d sera potentiellement égal à 1. Étapes suivantes Différencier les données et réanalyser les graphes ACF et PACF. Tester un modèle ARIMA(p, d, q) en commençant par p=1, d=1, q=0

En conclusion, les résultats indiquent que la série temporelle testée est non-stationnaire. Pour rendre la série stationnaire, vous pouvez envisager de différencier les données ou d'appliquer une transformation appropriée.

```
In [375... #transform the closing price to make it stationary

# Log transformation
df['close_log'] = np.log(df['Close'])

# Differencing
df['close_log_diff'] = df['close_log'].diff()

# Remove infinite values and NaNs
df['close_log_diff'].replace([np.inf, -np.inf], np.nan, inplace=True)
df['close_log_diff'].dropna(inplace=True)
```

Transformation logarithmique (log) :

La transformation logarithmique est utilisée pour stabiliser la variance de la série temporelle. Cela est particulièrement utile si la série présente une tendance exponentielle ou si les variations augmentent avec le temps.

Différenciation (diff()) :

La différenciation est utilisée pour éliminer les tendances et rendre la série temporelle stationnaire en termes de moyenne. En prenant la différence entre les valeurs consécutives,

les tendances linéaires sont éliminées.

- **diff()** : Cette méthode calcule la différence entre chaque élément et l'élément précédent dans la série temporelle. Elle est souvent utilisée pour transformer une série non stationnaire en une série stationnaire en éliminant les tendances.

En résumé, appliquer la transformation logarithmique suivie de la différenciation peut aider à rendre une série temporelle stationnaire en stabilisant la variance et en éliminant les tendances.

In [415... `df.head()`

Out[415...

	Close	MA20	MA50	Close_scaled	close_log	close_log_diff
Date						
2021-12-14	174.330002	NaN	NaN	0.400536	5.160950	NaN
2021-12-15	179.300003	NaN	NaN	0.440907	5.189060	0.028110
2021-12-16	172.259995	NaN	NaN	0.383722	5.149005	-0.040055
2021-12-17	171.139999	NaN	NaN	0.374624	5.142482	-0.006523
2021-12-20	169.750000	NaN	NaN	0.363334	5.134327	-0.008155


In [376... *#retest the stationarity of the closing price*

```
def test_stationarity(timeseries):  
  
    # Transformation to make the data stationary  
    timeseries_log = np.log(timeseries)  
    timeseries_log_diff = timeseries_log.diff().dropna()  
  
    # Remove infinite values and NaNs  
    timeseries_log_diff.replace([np.inf, -np.inf], np.nan, inplace=True)  
    timeseries_log_diff.dropna(inplace=True)  
  
    # Determining rolling statistics  
    rolmean = timeseries_log_diff.rolling(window=12).mean()  
    rolstd = timeseries_log_diff.rolling(window=12).std()  
  
    # Plot rolling statistics  
    plt.figure(figsize=(6, 4))  
    plt.plot(timeseries_log_diff, color='blue', label='Original')  
    plt.plot(rolmean, color='red', label='Rolling Mean')  
    plt.plot(rolstd, color='black', label='Rolling Std')  
    plt.legend(loc='best')  
    plt.title('Rolling Mean and Standard Deviation')  
    plt.xlabel('Date')  
    plt.ylabel('Log-Difference')  
    plt.show()  
  
    # Perform ADF test  
    adft = adfuller(timeseries_log_diff, autolag='AIC')
```

```
# Output the ADF test results
output = pd.Series(adft[0:4], index=['Test Statistic', 'p-value', 'Number of La
for key, value in adft[4].items():
    output[f'Critical Value ({key})'] = value
print(output)

test_stationarity(df['close_log_diff'])
```


C:\Users\darck\AppData\Roaming\Python\Python311\site-packages\pandas\core\arraylike.py:396: RuntimeWarning: divide by zero encountered in log
 result = getattr(ufunc, method)(*inputs, **kwargs)
C:\Users\darck\AppData\Roaming\Python\Python311\site-packages\pandas\core\arraylike.py:396: RuntimeWarning: invalid value encountered in log
 result = getattr(ufunc, method)(*inputs, **kwargs)




Test Statistic	-1.736206e+01
p-value	5.196850e-30
Number of Lags Used	0.000000e+00
Number of Observations Used	2.120000e+02
Critical Value (1%)	-3.461578e+00
Critical Value (5%)	-2.875272e+00
Critical Value (10%)	-2.574089e+00
dtype:	float64

En conclusion, les résultats indiquent que la série temporelle différenciée en logarithme est stationnaire. Cela signifie que les propriétés statistiques de la série (telles que la moyenne et la variance) ne changent pas au fil du temps, ce qui est une condition préalable importante pour nos modèles de prévision de séries temporelles.

```
In [377... #acf and pacf of the transformed closing price  
plot_acf(df['close_log_diff'], lags=20)  
plt.show()  
  
plot_pacf(df['close_log_diff'], lags=20)  
plt.show()
```





BUILD MODEL ARIMA

Commençant avec un modèle ARIMA(p, d, q) avec d=1 (différenciation pour stationnarité), p=1 ou 2 (basé sur PACF), et q=0 (faible influence MA suggérée).

```
In [378... #replace the infinite values with NaNs
df['close_log_diff'].replace([np.inf, -np.inf], np.nan, inplace=True)

#drop the NaNs
df['close_log_diff'].dropna(inplace=True)
```

```
In [379... #split the data into training and testing sets
train_size = int(len(df) * 0.8)
train_log_diff, test_log_diff = df['close_log_diff'][:train_size], df['close_log_diff'][train_size:]

train_log_diff.head()
```

```
Out[379... Date
2021-12-14      NaN
2021-12-15    0.028110
2021-12-16   -0.040055
2021-12-17   -0.006523
2021-12-20   -0.008155
Name: close_log_diff, dtype: float64
```

```
In [380... #drop the NaNs
train_log_diff.dropna(inplace=True)
```

```
In [381... #check for missing values
train_log_diff.isna().sum()
```

```
Out[381... 0
```

```
In [382... #build the ARIMA model with the training data and make predictions on the testing d

# Ensure the data is in the correct format
train_log_diff = train_log_diff.dropna()

# Build the ARIMA model
model = ARIMA(train_log_diff, order=(1, 1, 1))
model_fit = model.fit()
print(model_fit.summary())

# Make predictions
predictions1 = model_fit.forecast(steps=len(test_log_diff))

# Plot the predictions
plt.figure(figsize=(10, 6))
plt.plot(train_log_diff, label='Training')
plt.plot(test_log_diff, label='Actual')
plt.plot(test_log_diff.index, predictions1, label='Predicted')
plt.legend()
plt.show()
```

```
C:\Users\darck\AppData\Roaming\Python\Python311\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
```

```
self._init_dates(dates, freq)
```

```
C:\Users\darck\AppData\Roaming\Python\Python311\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
```

```
self._init_dates(dates, freq)
```

```
C:\Users\darck\AppData\Roaming\Python\Python311\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.
```

```
self._init_dates(dates, freq)
```

```
C:\Users\darck\AppData\Roaming\Python\Python311\site-packages\statsmodels\base\model.py:607: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals
```

```
warnings.warn("Maximum Likelihood optimization failed to "
```

```
C:\Users\darck\AppData\Roaming\Python\Python311\site-packages\statsmodels\tsa\base\tsa_model.py:836: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.
```

```
return get_prediction_index(
```

```
C:\Users\darck\AppData\Roaming\Python\Python311\site-packages\statsmodels\tsa\base\tsa_model.py:836: FutureWarning: No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.
```

```
return get_prediction_index(
```

SARIMAX Results

```

=====
Dep. Variable:      close_log_diff    No. Observations:      603
Model:              ARIMA(1, 1, 1)    Log Likelihood          1567.717
Date:              Sat, 14 Dec 2024    AIC                    -3129.433
Time:              16:11:14           BIC                    -3116.233
Sample:            0                  HQIC                   -3124.295
                  - 603
Covariance Type:    opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.0085	0.036	0.235	0.814	-0.062	0.079
ma.L1	-0.9999	0.648	-1.542	0.123	-2.271	0.271
sigma2	0.0003	0.000	1.544	0.123	-8.54e-05	0.001


```

=====
Ljung-Box (L1) (Q):      0.01    Jarque-Bera (JB):      104.24
Prob(Q):                 0.93    Prob(JB):           0.00
Heteroskedasticity (H):  0.39    Skew:               0.15
Prob(H) (two-sided):     0.00    Kurtosis:           5.02
=====

```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).



In [383...

```

mse = mean_squared_error(test_log_diff, predictions1)
mse

#calculate the accuracy of the model
accuracy = 100 - (mse * 100)
print(f'Model Accuracy: {accuracy}%', '\n', f'mean_squared_error : {mse}')
```

```

Model Accuracy: 99.98017605931706%
mean_squared_error : 0.0001982394068293454
```


BUILD MODEL SARIMA

In [384... *# Build the SARIMA model with the training data and make predictions on the testing*

```
# Build the SARIMA model
model = SARIMAX(train_log_diff, order=(1, 1, 1), seasonal_order=(1, 1, 1, 12))
model_fit = model.fit()
print(model_fit.summary())

# Make predictions
predictions2 = model_fit.forecast(steps=len(test_log_diff))

# Plot the predictions
plt.figure(figsize=(10, 6))
plt.plot(train_log_diff, label='Training')
plt.plot(test_log_diff, label='Actual')
plt.plot(test_log_diff.index, predictions2, label='Predicted')
plt.legend()
plt.show()
```

C:\Users\darck\AppData\Roaming\Python\Python311\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
self._init_dates(dates, freq)
```

C:\Users\darck\AppData\Roaming\Python\Python311\site-packages\statsmodels\tsa\base\tsa_model.py:473: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

```
self._init_dates(dates, freq)
```

C:\Users\darck\AppData\Roaming\Python\Python311\site-packages\statsmodels\base\model.py:607: ConvergenceWarning: Maximum Likelihood optimization failed to converge. Check mle_retvals

```
warnings.warn("Maximum Likelihood optimization failed to "
```

C:\Users\darck\AppData\Roaming\Python\Python311\site-packages\statsmodels\tsa\base\tsa_model.py:836: ValueWarning: No supported index is available. Prediction results will be given with an integer index beginning at `start`.

```
return get_prediction_index(
```

C:\Users\darck\AppData\Roaming\Python\Python311\site-packages\statsmodels\tsa\base\tsa_model.py:836: FutureWarning: No supported index is available. In the next version, calling this method in a model without a supported index will result in an exception.

```
return get_prediction_index(
```

SARIMAX Results

Dep. Variable: close_log_diff No. Observations: 603
Model: SARIMAX(1, 1, 1)x(1, 1, 1, 12) Log Likelihood 15
06.908
Date: Sat, 14 Dec 2024 AIC -30
03.817
Time: 16:11:15 BIC -29
81.916
Sample: 0 HQIC -29
95.285


Covariance Type: opg

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	0.0051	0.037	0.138	0.890	-0.067	0.078
ma.L1	-0.9973	0.035	-28.420	0.000	-1.066	-0.929
ar.S.L12	-0.0442	0.039	-1.148	0.251	-0.120	0.031
ma.S.L12	-0.9992	1.283	-0.779	0.436	-3.513	1.515
sigma2	0.0003	0.000	0.791	0.429	-0.000	0.001

Ljung-Box (L1) (Q): 0.00 Jarque-Bera (JB): 86.19
Prob(Q): 0.96 Prob(JB): 0.00
Heteroskedasticity (H): 0.41 Skew: 0.18
Prob(H) (two-sided): 0.00 Kurtosis: 4.84

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).



```
In [385... #calculate the mean squared error of the predictions
mse = mean_squared_error(test_log_diff, predictions2)

#calculate the accuracy of the model
accuracy = 100 - (mse * 100)
print(f"accuracy = " ,accuracy,"error = " , mse)
```

accuracy = 99.98053953763929 error = 0.00019460462360715368

BUILD MODEL LSTM

```
In [386... #display the data
df.head()
```

```
Out[386...      Close  MA20  MA50  Close_scaled  close_log  close_log_diff

Date
2021-12-14  174.330002   NaN   NaN      0.400536  5.160950          NaN
2021-12-15  179.300003   NaN   NaN      0.440907  5.189060      0.028110
2021-12-16  172.259995   NaN   NaN      0.383722  5.149005     -0.040055
2021-12-17  171.139999   NaN   NaN      0.374624  5.142482     -0.006523
2021-12-20  169.750000   NaN   NaN      0.363334  5.134327     -0.008155
```

```
In [387... # Create a new dataframe with only the 'Close column
data1 = df.filter(['Close'])
# Convert the dataframe to a numpy array
dataset = data1.values
# Get the number of rows to train the model on
training_data_len = int(np.ceil( len(dataset) * .95 ))

training_data_len
```

Out[387... 718

```
In [388... #reshape the Close_scaled column to a 2D array
scaled_data = df['Close_scaled'].values.reshape(-1,1)
```

```
In [390... # Create the training data set
# Create the scaled training data set
train_data = scaled_data[0:int(training_data_len), :]
# Split the data into x_train and y_train data sets
x_train = []
y_train = []

for i in range(60, len(train_data)):
    x_train.append(train_data[i-60:i, 0])
    y_train.append(train_data[i, 0])
    if i<= 61:
```

```

        print(x_train)
        print(y_train)
        print()

# Convert the x_train and y_train to numpy arrays
x_train, y_train = np.array(x_train), np.array(y_train)

# Reshape the data
x_train = np.reshape(x_train, (x_train.shape[0], x_train.shape[1], 1))

[array([0.40053612, 0.44090653, 0.38372183, 0.37462432, 0.36333361,
        0.38965158, 0.41117699, 0.41637559, 0.44927302, 0.44082522,
        0.44155637, 0.43197138, 0.42685409, 0.46291929, 0.44415561,
        0.40532855, 0.38160994, 0.38299081, 0.3831533 , 0.40662823,
        0.41028348, 0.3831533 , 0.39030142, 0.36373977, 0.33474126,
        0.32077 , 0.30371216, 0.29729507, 0.28234912, 0.28161809,
        0.27780036, 0.36804486, 0.40419136, 0.4028105 , 0.41280153,
        0.38892043, 0.38477784, 0.37884822, 0.40459753, 0.41637559,
        0.38258464, 0.35431728, 0.3562668 , 0.38802691, 0.38607752,
        0.3562668 , 0.34343273, 0.31922677, 0.2847048 , 0.30639271,
        0.32353186, 0.32572493, 0.31012913, 0.3374218 , 0.33474126,
        0.30988546, 0.2784502 , 0.26334176, 0.30809843, 0.27211441])],
[0.24132886927125385]

[array([0.40053612, 0.44090653, 0.38372183, 0.37462432, 0.36333361,
        0.38965158, 0.41117699, 0.41637559, 0.44927302, 0.44082522,
        0.44155637, 0.43197138, 0.42685409, 0.46291929, 0.44415561,
        0.40532855, 0.38160994, 0.38299081, 0.3831533 , 0.40662823,
        0.41028348, 0.3831533 , 0.39030142, 0.36373977, 0.33474126,
        0.32077 , 0.30371216, 0.29729507, 0.28234912, 0.28161809,
        0.27780036, 0.36804486, 0.40419136, 0.4028105 , 0.41280153,
        0.38892043, 0.38477784, 0.37884822, 0.40459753, 0.41637559,
        0.38258464, 0.35431728, 0.3562668 , 0.38802691, 0.38607752,
        0.3562668 , 0.34343273, 0.31922677, 0.2847048 , 0.30639271,
        0.32353186, 0.32572493, 0.31012913, 0.3374218 , 0.33474126,
        0.30988546, 0.2784502 , 0.26334176, 0.30809843, 0.27211441)], array([0.440906
53, 0.38372183, 0.37462432, 0.36333361, 0.38965158,
        0.41117699, 0.41637559, 0.44927302, 0.44082522, 0.44155637,
        0.43197138, 0.42685409, 0.46291929, 0.44415561, 0.40532855,
        0.38160994, 0.38299081, 0.3831533 , 0.40662823, 0.41028348,
        0.3831533 , 0.39030142, 0.36373977, 0.33474126, 0.32077 ,
        0.30371216, 0.29729507, 0.28234912, 0.28161809, 0.27780036,
        0.36804486, 0.40419136, 0.4028105 , 0.41280153, 0.38892043,
        0.38477784, 0.37884822, 0.40459753, 0.41637559, 0.38258464,
        0.35431728, 0.3562668 , 0.38802691, 0.38607752, 0.3562668 ,
        0.34343273, 0.31922677, 0.2847048 , 0.30639271, 0.32353186,
        0.32572493, 0.31012913, 0.3374218 , 0.33474126, 0.30988546,
        0.2784502 , 0.26334176, 0.30809843, 0.27211441, 0.24132887])],
[0.24132886927125385, 0.20794408870696612]

```

In [409... **from** keras.layers **import** Dropout

```

# Build the improved LSTM model
model = Sequential()
model.add(LSTM(256, return_sequences=True, input_shape=(x_train.shape[1], 1)))

```

```

model.add(Dropout(0.2))
model.add(LSTM(128, return_sequences=True))
model.add(Dropout(0.2))
model.add(LSTM(64, return_sequences=False))
model.add(Dropout(0.2))
model.add(Dense(50))
model.add(Dense(1))

# Compile the model
model.compile(optimizer='adam', loss='mean_squared_error')

# Train the model
model.fit(x_train, y_train, batch_size=32, epochs=100)

```





























Epoch 1/100





























C:\Users\darck\AppData\Roaming\Python\Python311\site-packages\keras\src\layers\rnn\rnn.py:204: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.





























```

    super().__init__(**kwargs)

```

21/21		5s	66ms/step	-	loss: 0.0579
Epoch 2/100					
21/21		1s	68ms/step	-	loss: 0.0062
Epoch 3/100					
21/21		1s	70ms/step	-	loss: 0.0051
Epoch 4/100					
21/21		1s	67ms/step	-	loss: 0.0043
Epoch 5/100					
21/21		1s	69ms/step	-	loss: 0.0048
Epoch 6/100					
21/21		1s	69ms/step	-	loss: 0.0047
Epoch 7/100					
21/21		2s	72ms/step	-	loss: 0.0033
Epoch 8/100					
21/21		1s	69ms/step	-	loss: 0.0037
Epoch 9/100					
21/21		1s	69ms/step	-	loss: 0.0034
Epoch 10/100					
21/21		1s	68ms/step	-	loss: 0.0038
Epoch 11/100					
21/21		2s	72ms/step	-	loss: 0.0045
Epoch 12/100					
21/21		1s	69ms/step	-	loss: 0.0025
Epoch 13/100					
21/21		1s	67ms/step	-	loss: 0.0030
Epoch 14/100					
21/21		1s	70ms/step	-	loss: 0.0027
Epoch 15/100					
21/21		1s	65ms/step	-	loss: 0.0025
Epoch 16/100					
21/21		1s	68ms/step	-	loss: 0.0022
Epoch 17/100					
21/21		1s	69ms/step	-	loss: 0.0023
Epoch 18/100					
21/21		1s	68ms/step	-	loss: 0.0024
Epoch 19/100					
21/21		1s	69ms/step	-	loss: 0.0023
Epoch 20/100					
21/21		2s	72ms/step	-	loss: 0.0022
Epoch 21/100					
21/21		2s	75ms/step	-	loss: 0.0026
Epoch 22/100					
21/21		1s	70ms/step	-	loss: 0.0021
Epoch 23/100					
21/21		1s	69ms/step	-	loss: 0.0020
Epoch 24/100					
21/21		2s	71ms/step	-	loss: 0.0024
Epoch 25/100					
21/21		1s	69ms/step	-	loss: 0.0023
Epoch 26/100					
21/21		2s	74ms/step	-	loss: 0.0019
Epoch 27/100					
21/21		2s	71ms/step	-	loss: 0.0016
Epoch 28/100					
21/21		2s	71ms/step	-	loss: 0.0021
Epoch 29/100					

21/21		1s 70ms/step	- loss: 0.0019
Epoch 30/100			
21/21		2s 71ms/step	- loss: 0.0025
Epoch 31/100			
21/21		1s 70ms/step	- loss: 0.0017
Epoch 32/100			
21/21		1s 70ms/step	- loss: 0.0020
Epoch 33/100			
21/21		2s 70ms/step	- loss: 0.0017
Epoch 34/100			
21/21		1s 71ms/step	- loss: 0.0015
Epoch 35/100			
21/21		2s 77ms/step	- loss: 0.0015
Epoch 36/100			
21/21		2s 73ms/step	- loss: 0.0017
Epoch 37/100			
21/21		2s 71ms/step	- loss: 0.0016
Epoch 38/100			
21/21		1s 70ms/step	- loss: 0.0018
Epoch 39/100			
21/21		2s 73ms/step	- loss: 0.0014
Epoch 40/100			
21/21		1s 65ms/step	- loss: 0.0017
Epoch 41/100			
21/21		1s 70ms/step	- loss: 0.0016
Epoch 42/100			
21/21		2s 71ms/step	- loss: 0.0017
Epoch 43/100			
21/21		2s 72ms/step	- loss: 0.0014
Epoch 44/100			
21/21		2s 71ms/step	- loss: 0.0015
Epoch 45/100			
21/21		2s 72ms/step	- loss: 0.0014
Epoch 46/100			
21/21		2s 71ms/step	- loss: 0.0013
Epoch 47/100			
21/21		2s 71ms/step	- loss: 0.0014
Epoch 48/100			
21/21		2s 72ms/step	- loss: 0.0014
Epoch 49/100			
21/21		2s 72ms/step	- loss: 0.0014
Epoch 50/100			
21/21		1s 69ms/step	- loss: 0.0018
Epoch 51/100			
21/21		2s 71ms/step	- loss: 0.0013
Epoch 52/100			
21/21		2s 73ms/step	- loss: 0.0014
Epoch 53/100			
21/21		1s 69ms/step	- loss: 0.0013
Epoch 54/100			
21/21		2s 74ms/step	- loss: 0.0014
Epoch 55/100			
21/21		2s 72ms/step	- loss: 0.0013
Epoch 56/100			
21/21		2s 73ms/step	- loss: 0.0013
Epoch 57/100			

21/21		1s 70ms/step	- loss: 0.0014
Epoch 58/100			
21/21		2s 72ms/step	- loss: 0.0013
Epoch 59/100			
21/21		1s 69ms/step	- loss: 0.0013
Epoch 60/100			
21/21		2s 71ms/step	- loss: 0.0012
Epoch 61/100			
21/21		2s 73ms/step	- loss: 0.0012
Epoch 62/100			
21/21		2s 72ms/step	- loss: 0.0014
Epoch 63/100			
21/21		2s 71ms/step	- loss: 0.0012
Epoch 64/100			
21/21		2s 71ms/step	- loss: 0.0013
Epoch 65/100			
21/21		2s 71ms/step	- loss: 0.0012
Epoch 66/100			
21/21		2s 72ms/step	- loss: 0.0013
Epoch 67/100			
21/21		2s 71ms/step	- loss: 0.0012
Epoch 68/100			
21/21		2s 72ms/step	- loss: 0.0012
Epoch 69/100			
21/21		1s 69ms/step	- loss: 0.0014
Epoch 70/100			
21/21		2s 72ms/step	- loss: 0.0013
Epoch 71/100			
21/21		1s 70ms/step	- loss: 0.0012
Epoch 72/100			
21/21		2s 71ms/step	- loss: 0.0012
Epoch 73/100			
21/21		2s 72ms/step	- loss: 0.0011
Epoch 74/100			
21/21		2s 72ms/step	- loss: 0.0012
Epoch 75/100			
21/21		1s 70ms/step	- loss: 0.0012
Epoch 76/100			
21/21		1s 70ms/step	- loss: 0.0014
Epoch 77/100			
21/21		2s 72ms/step	- loss: 0.0013
Epoch 78/100			
21/21		1s 67ms/step	- loss: 0.0011
Epoch 79/100			
21/21		2s 72ms/step	- loss: 0.0012
Epoch 80/100			
21/21		2s 83ms/step	- loss: 0.0010
Epoch 81/100			
21/21		2s 78ms/step	- loss: 0.0013
Epoch 82/100			
21/21		2s 79ms/step	- loss: 0.0011
Epoch 83/100			
21/21		2s 78ms/step	- loss: 0.0013
Epoch 84/100			
21/21		2s 80ms/step	- loss: 0.0011
Epoch 85/100			


```

21/21 ————— 2s 79ms/step - loss: 0.0013
Epoch 86/100
21/21 ————— 2s 80ms/step - loss: 0.0013
Epoch 87/100
21/21 ————— 2s 79ms/step - loss: 0.0012
Epoch 88/100
21/21 ————— 2s 78ms/step - loss: 9.7751e-04
Epoch 89/100
21/21 ————— 2s 79ms/step - loss: 0.0011
Epoch 90/100
21/21 ————— 2s 77ms/step - loss: 0.0011
Epoch 91/100
21/21 ————— 2s 77ms/step - loss: 0.0013
Epoch 92/100
21/21 ————— 2s 78ms/step - loss: 0.0014
Epoch 93/100
21/21 ————— 2s 77ms/step - loss: 0.0014
Epoch 94/100
21/21 ————— 2s 73ms/step - loss: 0.0011
Epoch 95/100
21/21 ————— 1s 70ms/step - loss: 0.0010
Epoch 96/100
21/21 ————— 2s 72ms/step - loss: 0.0011
Epoch 97/100
21/21 ————— 1s 69ms/step - loss: 0.0013
Epoch 98/100
21/21 ————— 2s 73ms/step - loss: 0.0012
Epoch 99/100
21/21 ————— 2s 71ms/step - loss: 0.0010
Epoch 100/100
21/21 ————— 2s 72ms/step - loss: 0.0012

```

Out[409... <keras.src.callbacks.history.History at 0x2f002088b90>

In [410... `model.summary()`

Model: "sequential_8"

Layer (type)	Output Shape	Param #
lstm_24 (LSTM)	(None, 60, 256)	264,192
dropout (Dropout)	(None, 60, 256)	0
lstm_25 (LSTM)	(None, 60, 128)	197,120
dropout_1 (Dropout)	(None, 60, 128)	0
lstm_26 (LSTM)	(None, 64)	49,408
dropout_2 (Dropout)	(None, 64)	0
dense_16 (Dense)	(None, 50)	3,250
dense_17 (Dense)	(None, 1)	51

Total params: 1,542,065 (5.88 MB)
Trainable params: 514,021 (1.96 MB)
Non-trainable params: 0 (0.00 B)
Optimizer params: 1,028,044 (3.92 MB)

```
In [411... # Create the testing data set
# Create a new array containing scaled values from index 1543 to 2002
test_data = scaled_data[training_data_len - 60: , :]
# Create the data sets x_test and y_test
x_test = []
y_test = dataset[training_data_len:, 0] # Select only the 'Close' column
for i in range(60, len(test_data)):
    x_test.append(test_data[i-60:i, 0])

# Convert the data to a numpy array
x_test = np.array(x_test)

# Reshape the data
x_test = np.reshape(x_test, (x_test.shape[0], x_test.shape[1], 1))

# Get the models predicted price values
predictions3 = model.predict(x_test)
predictions3 = scaler.inverse_transform(predictions3)

# Get the root mean squared error (RMSE)
rmse = np.sqrt(np.mean(((predictions3 - y_test.reshape(-1, 1)) ** 2)))
rmse

#calculate the mean squared error of the predictions
mse = mean_squared_error(y_test, predictions3)
mse
```

2/2 ————— 2s 357ms/step

Out[411... 9.220745615642267


```
In [412... # Plot the data
train = data1[:training_data_len]
valid = data1[training_data_len:]
valid['Predictions'] = predictions3

# Visualize the data
plt.figure(figsize=(16,6))
plt.title('Model LSTM')
plt.xlabel('Date', fontsize=18)
plt.ylabel('Close Price USD ($)', fontsize=18)
plt.plot(train['Close'])
plt.plot(valid[['Close', 'Predictions']])
plt.legend(['Train', 'Val', 'Predictions'], loc='lower right')
plt.show()

# Show the accurate
```

C:\Users\darck\AppData\Local\Temp\ipykernel_26764\2554651337.py:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
valid['Predictions'] = predictions3



In [413...

```
# Calculer les métriques d'évaluation pour LSTM
mae = mean_absolute_error(y_test, predictions3)
print(f'Mean Absolute Error: {mae}')

mse = mean_squared_error(y_test, predictions3)
print(f'Mean Squared Error: {mse}')
```

Mean Absolute Error: 2.6506764179951436

Mean Squared Error: 9.220745615642267

In [414...

```
valid.head()
```

Out[414...

	Close	Predictions
Date		
2024-10-23	230.759995	234.465393
2024-10-24	230.570007	230.687195
2024-10-25	231.410004	229.500397
2024-10-28	233.399994	230.136993
2024-10-29	233.669998	231.900864